

Data Structure And Algorithms - Coursework

January 2018

Contents

1	Analysis of the problem	4
2	Proposed Data Structures	4
2.1	Year Map ADT	6
2.1.1	Lookup table	6
2.2	Annual Readings	6
2.2.1	Data Fields	6
2.2.2	Methods	7
2.2.3	PsuedoCode	7
2.3	Month Map ADT	7
2.3.1	Lookup Table	8
2.4	Days Array	9
2.5	Sensor Map ADT	9
2.5.1	Lookup Table	10
2.6	Daily Readings	10
2.6.1	Data fields	12
2.6.2	Methods	12
2.6.3	Pseudo Code	12
3	Algorithms	14
3.1	The maximum wind speed of a specified month and year	14
3.1.1	Pseudo Code	14
3.2	The median wind speed of a specified year	14
3.2.1	Pseudo Code	15
3.3	Average wind speed for each month of a specified year in the order of month	15
3.3.1	Pseudo Code	15
3.4	Total solar radiation for each month of a specified year in a descend- ing order of the solar radiation	16
3.4.1	Pseudo Code	16
3.5	Given a date, show the times for the highest solar radiation for that date, including duplicates, displayed in reverse chronological order .	16
3.5.1	Pseudo Code	17
4	Conclusion	17
4.1	Space and time requirements	17
4.1.1	The maximum wind speed of a specified month and year . .	17
4.1.2	The median wind speed of a specified year	18

4.1.3	Average wind speed for each month of a specified year in the order of month	18
4.1.4	Total solar radiation for each month of a specified year in a descending order of the solar radiation	18
4.1.5	Given a date, show the times for the highest solar radiation for that date, including duplicates, displayed in reverse chronological order	18

1 Analysis of the problem

The goal of the exercise is to store 10 years worth of 5 minute interval sensor readings in such a way that lookups can be done efficiently. One main characteristic of this data is that the storage requirements are constant. The client has 10 years of data and every year has 12 months. The number of days in every month is constant and so is the total number of individual sensor readings per day, which is 288. ie a reading every 5 minutes for a total 1440 minutes ($60 * 24$).

A common requirement in the problem set is to get the sensor readings on a specific date. In some cases we also need the timestamp associated with a reading. Considering the size(1 million readings per sensor), the data needs to be organized in a way to do lookups with maximum efficiency avoiding sequential access as much possible. Another distinct requirement is to find the median Wind Speed in a given year. This requires sorting 105,120 values and we need to use an appropriate data structure to do this along with insert operation to increase efficiency.

Each input file has an year worth of data in chronological order. The date together with timestamp per line is unique within and across files. The lines are processed one by one and inserted into the proposed data structure(s). Values like maximum, average and total are calculated along with insert to avoid redundant linear parsing later. The data structure should support this using custom methods. The fact that the data set may contain duplicates has to be considered while searching and sorting. Hence the search and sort algorithms must support duplicates.

2 Proposed Data Structures

We use a combination of Arrays, Map ADTs and Custom ADTs [1] to store the entire data. Let's assume that there is a wrapper function to process the input files and read one line at a time. It then separates the time stamp and sensor readings and insert them to the proposed data structures using various insertRecord() methods in the abstract data types.

Wind Speed reading for a year is stored twice. One copy in a combination of a minheap and a max heap in the AnnualReadings ADT and another copy in the specific DailyReadings ADT. By splitting the wind speed readings into a maxheap and minheap of equal node count, we ensure that the middle values are always the root nodes of the respective heaps. We are sacrificing space complexity here

to find the median with a constant time complexity. ie To find the median all we need to do is to get the root nodes of the heaps and find the average value.

Since the keys of the input data (Year, Month, Day and Sensor Type) are pre-known and unique a *Map Abstract Data Type* is one of the best choices to store the data. A map lets us do the lookups with constant time complexity, ie $O(1)$.

For example to get the data for Wind Speeds on 2nd February 2014, a single lookup over a combination of map abstract data types fetches the required daily reading data structure.

`READINGS[2014].monthHashmap[FEB][2][S] -> DailyReadings`

The below diagram shows a visualization of all the data structures included.

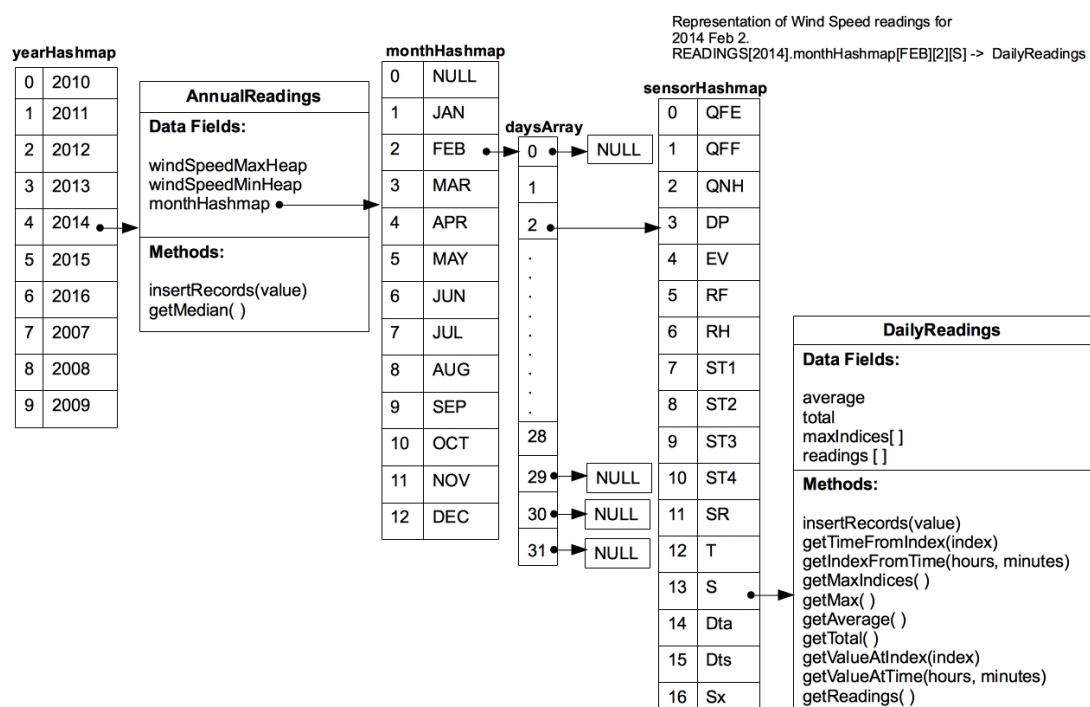


Figure 1: Wind Speeds on 2nd February 2014

Individual data structures shown in Figure 1 are discussed in detail below.

2.1 Year Map ADT

Data for a year is stored in the *yearHashmap*. Since the data set is smaller (10 entries), a simple modular hashing function with a linear collision resolution can be used on an array of 10 elements.

$$h(k) = k \bmod m$$

where $k = \text{year}$ and $m = 10$.

2.1.1 Lookup table

Key	Index
2007	7
2008	8
2009	9
2010	0
2011	1
2012	2
2013	3
2014	4
2015	5
2016	6

Element of yearHashmap is an object of AnnualReadings abstract data type.

`yearHashmap[2014] -> AnnualReadings`

2.2 Annual Readings

AnnualReadings custom ADT has the `windSpeedMinHeap`, `windSpeedMaxHeap` (used in combination to find the median) and the `monthHashmap` data members. It also has two methods. One to insert records to the heaps and another one to get the median wind speed. Only the Wind Speed sensor readings are inserted to the heaps in this ADT. Same readings are also stored in their corresponding *DailyReadings* ADT which is discussed in detail later in this section.

AnnualReadings ADT has the following blueprint

2.2.1 Data Fields

- `windSpeedMinHeap`

- windSpeedMaxHeap
- monthHashmap

2.2.2 Methods

- Insert new wind speed reading
- Get median wind speed

2.2.3 PsuedoCode

Algorithm 2.1: Blueprint of annualReadings ADT

```

1  CLASS AnnualReadings() :
2      SET windSpeedMinHeap
3      SET windSpeedMaxHeap
4      SET monthHashmap
5
6      FUNCTION insertRecord(value) :
7
8          // Insert value to the correct heap
9          IF windSpeedMaxHeap ROOT IS NULL:
10             INSERT value TO windSpeedMaxHeap
11          ELSE:
12             IF value > windSpeedMaxHeap ROOT:
13                 INSERT value to windSpeedMinHeap
14             ELSE:
15                 INSERT value to windSpeedMaxHeap
16             ENDIF
17          ENDIF
18
19          // Check the total number of nodes and balance the heaps if necessary
20          IF COUNT(windSpeedMaxHeap) + COUNT(windSpeedMinHeap) IS EVEN:
21             WHILE (COUNT(windSpeedMaxHeap) != COUNT(windSpeedMinHeap)) :
22                 IF COUNT(windSpeedMaxHeap) > COUNT(windSpeedMinHeap) :
23                     MOVE ROOT FROM windSpeedMaxHeap TO windSpeedMinHeap
24                 ELSE:
25                     MOVE ROOT FROM windSpeedMinHeap TO windSpeedMaxHeap
26                 ENDIF
27             ENDWHILE
28          ENDIF
29      ENDFUNCTION
30
31      FUNCTION getMedianWindSpeed() :
32          // Return the average of root nodes of min and max heaps
33          RETURN ((windSpeedMinHeap ROOT) + (windSpeedMaxHeap ROOT)) / 2
34      ENDFUNCTION
35  ENDCLASS

```

2.3 Month Map ADT

The *monthHashmap* uses a custom hashing function which accepts the 3 letter string representation of month name and output the numeric value of the month.

The output of this hashing function is always unique, hence there is no need for a collision resolution function. The size of the array used by this hashmap is 13 and index 0 is pointed to NULL. So that the actual values begins from 1 and goes up-to 12.

$h(\text{JAN}) = 1$
 $h(\text{FEB}) = 2$
 $h(\text{MAR}) = 3$
 $h(\text{APR}) = 4$
 $h(\text{MAY}) = 5$
 $h(\text{JUN}) = 6$
 $h(\text{JUL}) = 7$
 $h(\text{AUG}) = 8$
 $h(\text{SEP}) = 9$
 $h(\text{OCT}) = 10$
 $h(\text{NOV}) = 11$
 $h(\text{DEC}) = 12$

2.3.1 Lookup Table

Key	Index
JAN	1
FEB	2
MAR	3
APR	4
MAY	5
JUN	6
JUL	7
AUG	8
SEP	9
OCT	10
NOV	11
DEC	12

Element of monthHashmap is a daysArray.

`monthHashmap[FEB] -> daysArray []`

2.4 Days Array

The *daysArray* is an array of *sensorHashmap* values. Use of a map ADT here will be redundant and the day number can be used as the index of the array without passing it through a hashing function. The size of the array is 32, indexed (0..31), with all values initialized to NULL.

`daysArray = [NULL, NULL, . . . , NULL]`

Initializing the array elements to NULL helps in handling the varying number of days across months. For example if the `daysArray[31]` is pointing to NULL, it can be safely assumed that the month has only 30 days. However while looping through the *daysArray*, we need to keep a separate counter variable to get the actual number of days.

Algorithm 2.2: Count number of days in a given *daysArray*.

```
1  SET counter = 0
2
3  FOR index IN 1 TO 31:
4    IF daysArray[index] NOT NULL:
5      SET counter = counter + 1
6    ENDIF
7  ENDFOR
```

Element of daysArray is a sensorHashmap.

2.5 Sensor Map ADT

The *sensorHashmap* uses a custom hashing function as well. It takes the sensor type as an input and returns the index.

`h(QFE) = 0`
`h(QFF) = 1`
`h(QNH) = 2`
`h(DP) = 3`
`h(EV) = 4`
`h(RF) = 5`
`h(RH) = 6`
`h(ST1) = 7`
`h(ST2) = 8`
`h(ST3) = 9`
`h(ST4) = 10`
`h(SR) = 11`
`h(T) = 12`
`h(S) = 13`
`h(Dta) = 14`

$h(Dts) = 15$

$h(Sx) = 16$

2.5.1 Lookup Table

Key	Value
QFE	0
QFF	1
QNH	2
DP	3
EV	4
RF	5
RH	6
ST1	7
ST2	8
ST3	9
ST4	10
SR	11
T	12
S	13
Dta	14
Dts	15
Sx	16

Element of a sensorHashmap is an object of DailyReadings ADT.

2.6 Daily Readings

The goal in the design of data set for daily readings is to do the following actions efficiently:

1. Insert new item
2. Find the max value(s) including duplicates along with the timestamp
3. Find average value
4. Find total

We need to use a custom Abstract Data Type with sensor readings, max, average and total as data members and a set of methods for the above calculations.

Daily sensor readings are inserted to an *array* data field in this abstract data structure. However we need some special considerations for storing the readings.

It is a requirement that we not only need the maximum values but also the timestamp associated with them. We know that readings are taken every 5 minutes and that makes a total of 288 readings per day. If we get the readings sorted by time and store them in an array of size 289 (ie index 0 - 288) starting from index 1, the *n*th item will be the reading at (*n* * 5)th minute.

In 24 hr format we usually represent time as 00:00 to 23:59. That is the 288th reading in our case should actually be the first reading of the next day. The below algorithm can be tweaked to make it more realistic and support this case. However, for the sake of simplicity, we are assuming that the first reading of the day is taken at 00:05 and the last reading at 24:00. With this assumption, given an index *n*, we can find the corresponding time as below.

```
totalMinutes = n * 5  
hours = |totalMinutes/60|  
minutes = |totalMinutes mod 60|
```

For example:

readings[15] is the reading at 75th minute or in other words, the time will be 01:15 (24 hr format). The calculation can be reversed to find the reading at a specific time as follows.

```
totalMinutes = (hours * 60) + minutes  
n = totalMinutes/5  
value = readings[n]
```

Average and total are calculated along with insert operation and are stored as floating point data fields. Max value(s) might have duplicates and needs to retain the timestamp information. Since we can easily find the value and time from the index of the item in the sensor readings array, the max values can be represented as an array of their indices in the sensor readings array.

Putting all this together, an abstract data structure with the following blue print can be used to store the daily readings.

2.6.1 Data fields

- Array of readings
- Array of indices of max readings
- Average
- Total

2.6.2 Methods

- Insert new reading
- Return array of readings
- Return max readings array
- Return max reading
- Return total
- Return average
- Return time given an index
- Return index given a time

2.6.3 Pseudo Code

Algorithm 2.3: Blueprint of dailyReadings ADT

```
1 CLASS DailyReadings():
2   SET maxIndices = []
3   SET average = 0.0
4   SET total = 0.0
5   SET readings = []
6
7   FUNCTION insertRecord(value):
8     PUSH value TO readings[]
9     SET lastIndex = (length of readings[]) - 1
10
11     IF (length of maxIndices[]) > 0:
12       // Get the value at current max index
13       SET currentMax = readings[maxIndices[0]]
14
15       IF readings[lastIndex] > currentMax:
16         // Re-initialize the maxIndices array with lastIndex as the only member
17         SET maxIndices = [lastIndex]
18       ELSEIF readings[lastIndex] == currentMax:
19         // Push lastIndex to maxIndices array if value is equal to the max value
20         (duplicate)
```

```

20     PUSH lastIndex TO maxIndices []
21     ENDIF
22
23     ELSE:
24         // If max indices array is empty add this one.
25         PUSH lastIndex TO maxIndices
26     ENDIF
27
28     SET total = total + value
29     SET average = total / (length of readings [])
30 ENDFUNCTION
31
32 FUNCTION getTimeFromIndex(index):
33     // Assuming that time is the string representation of HH:MM in 24 hours
34     format
35
36     SET totalMinutes = index * 5
37     SET hour = INTEGER(totalMinutes / 60)
38     SET minutes = INTEGER(totalMinutes % 60)
39
40     RETURN hour + ":" + minutes
41 ENDFUNCTION
42
43 FUNCTION getIndexFromTime(hours, minutes):
44     SET totalMinutes = (hours * 60) + minutes
45     SET index = totalMinutes / 5
46     RETURN index
47 ENDFUNCTION
48
49 FUNCTION getMaxIndices():
50     RETURN maxIndices
51 ENDFUNCTION
52
53 FUNCTION getMax():
54     // Return the numeric max value (no duplicates)
55     IF (length of maxIndices) > 0:
56         RETURN readings[maxIndices[0]]
57     ELSE:
58         RETURN NULL
59     ENDIF
60 ENDFUNCTION
61
62 FUNCTION getAverage():
63     RETURN average
64 ENDFUNCTION
65
66 FUNCTION getTotal():
67     RETURN total
68 ENDFUNCTION
69
70 FUNCTION getValueAtIndex(index):
71     RETURN readings[index]
72 ENDFUNCTION
73
74 FUNCTION getValueAtTime(hours, minutes):
75     SET index = getIndexFromTime(hours, minutes)
76     RETURN getValueAtIndex(index)
77 ENDFUNCTION
78
79 FUNCTION getReadings():
80     RETURN readings []
ENDFUNCTION

```

3 Algorithms

The pseudo codes [2] in this section calls the methods in the *DailyReadings* and *AnnualReadings* abstract data types. This is highlighted using in-line comments.

3.1 The maximum wind speed of a specified month and year

Input: Integer year, String month

Output: Float windSpeed

Complexity: Constant, $O(1)$

3.1.1 Pseudo Code

Algorithm 3.1: Find maximum wind speed in a given month and year

```

1  FUNCTION getMaxWindSpeed(year , month):
2      SET max = 0
3
4      FOR index IN 1 TO 31:
5          SET day = READINGS[year].monthHashMap[month][index]
6
7          IF day NOT NULL:
8              SET dayMax = day[S].getMax()
9
10             IF dayMax > max:
11                 SET max = dayMax
12             ENDIF
13
14         ENDIF
15     ENDFOR
16
17     RETURN max
18 ENDFUNCTION

```

3.2 The median wind speed of a specified year

Input: Integer year

Output: Float median

Complexity: $O(1)$

3.2.1 Pseudo Code

Algorithm 3.2: Find median wind speed in a given year

```
1 FUNCTION getMedian(year):  
2   // Call the getMedianWindSpeed method in AnnualReadings ADT  
3   SET median = READINGS[year].getMedianWindSpeed()  
4   RETURN median  
5 ENDFUNCTION
```

3.3 Average wind speed for each month of a specified year in the order of month

Input: Integer year

Output: Array of Float averageWindSpeed sorted in the order of month

Complexity: Constant, $O(1)$

3.3.1 Pseudo Code

Algorithm 3.3: Find average wind speed per month for a given year and return values in the order of month

```
1 FUNCTION getAverageWindSpeed(year):  
2   SET averageWindSpeed = []  
3  
4   FOR month IN (JAN FEB MAR APR MAY JUN JUL AUG SEP OCT NOV DEC):  
5     SET monthAverageTotal = 0  
6     SET numberOfDays = 0  
7  
8     FOR index 1 TO 31:  
9       SET day = READINGS[year].monthHashMap[month][index]  
10  
11      IF day NOT NULL:  
12        SET numberOfDays = numberOfDays + 1  
13        // getAverage() returns the average value for the day  
14        SET monthAverageTotal = monthAverageTotal + day[S].getAverage()  
15      ENDIF  
16    ENDFOR  
17  
18    // average of averages == average (since weight is same; 288 readings per  
19    // day)  
19    SET monthAverage = monthAverageTotal / numberOfDays  
20    PUSH monthAverage TO averageWindSpeed[]  
21  ENDFOR  
22  
23  RETURN averageWindSpeed  
24 ENDFUNCTION
```

3.4 Total solar radiation for each month of a specified year in a descending order of the solar radiation

Input: Integer year

Output: Array of Float totalSolarRadiation sorted in descending order

Complexity: Linear, $O(n)$.

3.4.1 Pseudo Code

Algorithm 3.4: Return total solar radiation of each month in a given year; in descending order.

```
1  FUNCTION getTotalSolarRadiation(year):
2      SET totalSolarRadiation = []
3
4      FOR month IN (JAN FEB MAR APR MAY JUN JUL AUG SEP OCT NOV DEC):
5          SET monthTotal = 0
6
7          FOR index IN 1 TO 31:
8              SET day = READINGS[year].monthHashMap[month][index]
9
10             IF day NOT NULL:
11                 // getTotal() returns the sum of readings per day
12                 SET monthTotal = monthTotal + day[SR].getTotal()
13             ENDIF
14         ENDFOR
15
16         // Insert monthTotal into correct position in the array
17         FOR index IN 0 TO 11:
18             IF totalSolarRadiation[index] IS NULL:
19                 SET totalSolarRadiation[index] = monthTotal
20                 BREAKFOR
21             ELSEIF monthTotal > totalSolarRadiation[index]:
22                 PUSH ELEMENTS FROM INDEX UPTO END ONE STEP TO THE RIGHT
23                 SET totalSolarRadiation[index] = monthTotal
24                 BREAKFOR
25             ENDIF
26         ENDFOR
27
28     ENDFOR
29
30     RETURN totalSolarRadiation
31 ENDFUNCTION
```

3.5 Given a date, show the times for the highest solar radiation for that date, including duplicates, displayed in reverse chronological order

Input: Integer Year, String Month, Integer Day

Output: String representation of times as HH:MM in 24 hour format in reverse chronological order

Complexity: Linear, $O(n)$ - where n is the number of duplicates.

3.5.1 Pseudo Code

Algorithm 3.5: Show the times for the highest solar radiation for a given date, including duplicates, displayed in reverse chronological order

```
1 FUNCTION getHighestSolarRadiationTimes(year, month, day):
2
3   SET srDailyReadingsInstance = READINGS[year].monthHashMap[month][day][SR]
4
5   // getMaxIndices() returns the array of indices of max values
6   SET maxReadingsIndices = srDailyReadingsInstance.getMaxIndices()
7
8   // getMaxIndices() method returns the array of indexes sorted in chronological
   order by default
9   FOR index IN (length of maxReadingsIndices - 1) TO 0:
10    SET maxReadingIndex = maxReadingsIndices[index]
11
12    // getTimeFromIndex() converts the index to HH:MM string representation
13    SET time = srDailyReadingsInstance.getTimeFromIndex(maxReadingIndex)
14    PRINT time
15  ENDFOR
16
17 ENDFUNCTION
```

4 Conclusion

Most of the client requirements can be done in constant time complexity because of the way *inserts* are handled and using map abstract data type. The insert operation has a linear complexity $O(n)$, where n is the number of readings. The operations like finding max, sum and average which themselves has a linear complexity are done along with the insert. Moreover, the readings are saved in a chronologically sorted array. This way, these operations are completely removed from the other algorithms which are called more frequently compared to insert. *For example, Finding the total solar radiation in a month may be called multiple times, but inserting the daily solar radiation values is only done once.* By segregating the data across various maps; lookups are also done in constant time rather than looping through the entire data one at a time.

4.1 Space and time requirements

4.1.1 The maximum wind speed of a specified month and year

Though the algorithm loops through the *daysArray* the relative time complexity can be considered constant because n is always between 28 and 31. Space complex-

ity is also constant as we compare the daily maximum value with overall maximum and the same variables are re-used across the loop iterations.

4.1.2 The median wind speed of a specified year

To find the median we are using two different heaps in the *AnnualReadings* ADT. The `insertRecords()` member function creates these heaps by splitting the readings into two halves and storing them in a max heap and a min heap. The roots of these heaps will be mid points of all the values. So to calculate the median we need to find the average of these roots. Though the heap creation has a time complexity of $O(n \log n)$, this is handled along with insert. Hence finding the median can be done with a constant complexity $O(1)$. Space complexity is $O(2n)$ (or linear after removing the constant) because the values are saved twice.

4.1.3 Average wind speed for each month of a specified year in the order of month

The *monthHashmap* is visited in the order of the months and monthly averages are pushed into the output array in the same order. Hence the output array has been sorted in the required sequence. Since the *DailyReadings* class object has the average of readings as a data member, getting it has a constant complexity too. The space requirement is linear depending on the number of months with available readings.

4.1.4 Total solar radiation for each month of a specified year in a descending order of the solar radiation

Getting the total values can be done in constant complexity, however the results are sorted in place using insertion sort. Hence the algorithm has a time complexity of $O(n)$ and space complexity of $O(n)$.

4.1.5 Given a date, show the times for the highest solar radiation for that date, including duplicates, displayed in reverse chronological order

The algorithm loops through the *maxIndices* array, which is a data member of *DailyReadings* class. The iterations depends on the number of duplicate values and has a linear complexity, $O(n)$. The *maxIndices* array is retrieved and stored before processing and the space complexity is also linear.

References

- [1] M. Goodrich, R. Tamassia, and M. Goldwasser. *Data Structures and Algorithms in Java*. 6th ed. 2014.
- [2] JD. *Pseudocode Standard*. URL: http://users.csc.calpoly.edu/~jdalbey/SWE/pdl_std.html (visited on 01/06/2018).